




In the previous
lecture ...

Markov Decision Process

A **Markov decision process** is a tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ where

- \mathcal{S} is a finite set of states.
- \mathcal{A} is a finite set of actions.
- \mathcal{P} is a state-action transition probability matrix,
 - $\mathcal{P}_{s,s'}^a = \mathbb{P}(S_{t+1} = s' \mid S_t = s, A_t = a)$.
- \mathcal{R} is a reward function
 - $\mathcal{R}_s^a = \mathbb{E}(R_{t+1} \mid S_t = s, A_t = a)$
- γ is a **discount factor**, $\gamma \in [0, 1]$.

Let's define an MPD for space invaders



Policies

A **policy** π is a distribution over actions given states,

$$\pi(a \mid s) := \mathbb{P}(A_t = a \mid S_t = s).$$

- A policy fully defines the behaviour of the agent.
- Policies are stationary and depend only on the current state, not the history.

Recursion equations

As with MRP's, we have similar recursive relations:


$$\begin{aligned}v_{\pi}(s) &= \mathbb{E}_{\pi}(R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s) \\q_{\pi}(s, a) &= \mathbb{E}_{\pi}(R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a)\end{aligned}$$

In particular, we can, for a given policy π , solve the Bellman equation.


Note also that:

$$v_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) q_{\pi}(s, a).$$

S1 =



S2 =




Which state has a higher value?


If we play a reasonable policy.



Deep Reinforcement Learning



Collective Robot Reinforcement Learning. Training Phase



Neural Architecture Search with Reinforcement Learning

Task 02004: You are looking for and it should serve **gastropub food**. You don't care about the **price range**. You want to know the **address**.

Info Desk : Hello , welcome to the Cambridge restaurant system . You can ask for restaurants by area, price range or food type . How may I help you ?
 Customer : I want a gastropub food
 Info Desk : There are 4 restaurants serving gastropub food, what price range do you want ?

Next turn

Customer : (Your response)

A Network-based End-to-End Trainable Task-oriented Dialogue System



Mastering the game of Go with deep neural networks and tree search

Motivation

Reinforcement learning prior to 2014:

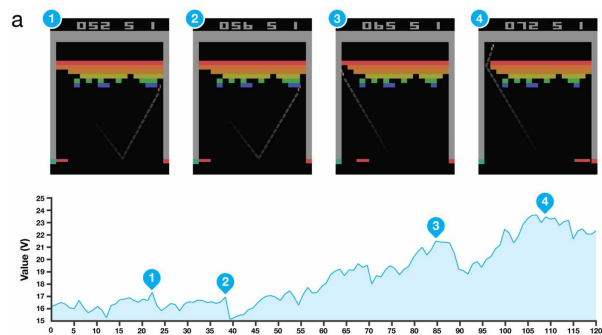
- handcrafted features, or fully observed, low-dimensional states
- achievements: a super-human Backgammon algorithm, a successful robot soccer agent



Reinforcement Learning for robot soccer - M.Riedmiller et al. (2009)

What we want:

- learn directly from high-dimensional sensory input, no feature engineering (end-to-end learning)
- one architecture that can excel at a variety of tasks



Human-level control through deep reinforcement learning - V.Mnih et al. (2014)


Solution

Deep Learning:

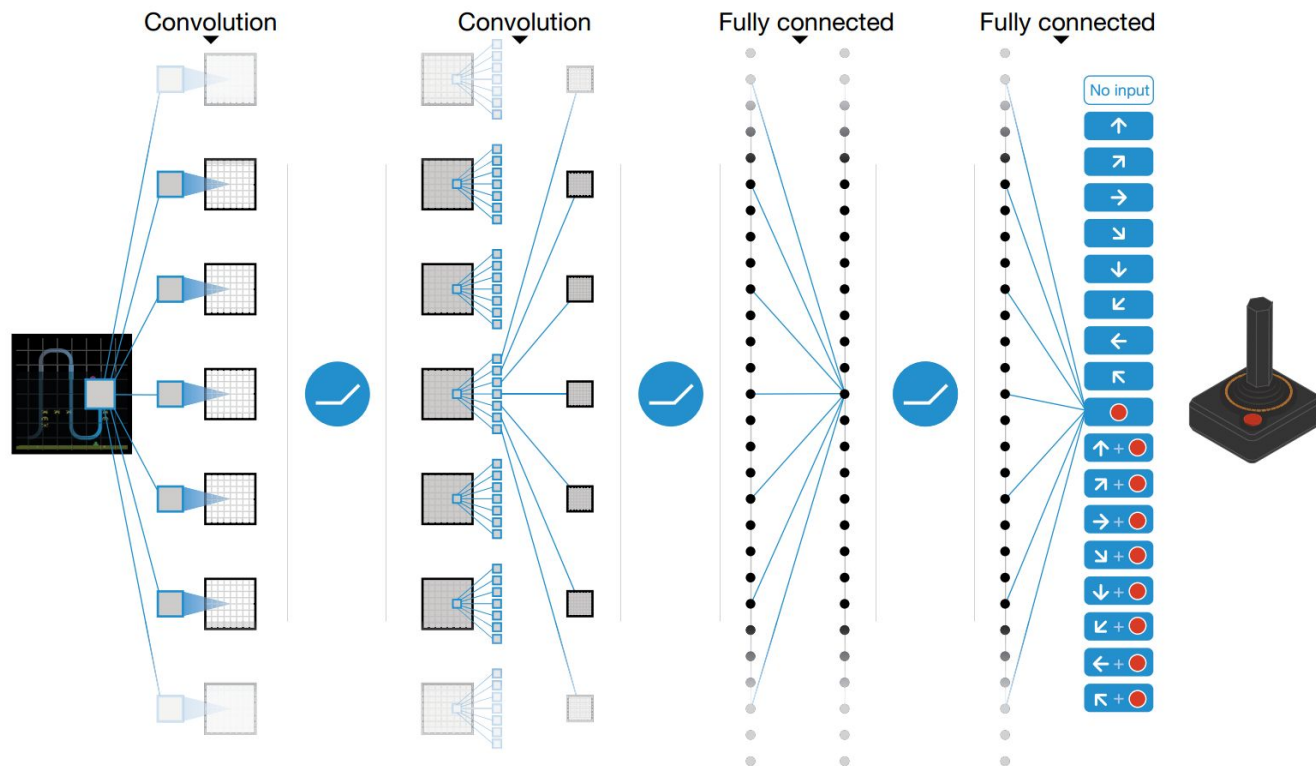
- Deep Neural Networks are capable of processing high-dimensional input and have been successfully applied to a wide range of problems
- we could use a DNN as a approximator for the action-value function (Q function)

Challenges:

Reinforcement learning is unstable when we use **nonlinear function approximators.**




Deep Q-Network





[DQN Breakout \(DeepMind 2016\)](#)



[DQN Space Invaders \(DeepMind 2016\)](#)


How Deep Q-Network works

- we are modeling the learning task as a Reinforcement Learning problem
- therefore, we are dealing with Markov Decision Processes
- **action-value function** for an optimal policy



$$Q^*(s,a) = \max_{\pi} \mathbb{E} [r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s, a_t = a, \pi]$$

- let's turn it into a **loss function** for our neural network:

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s',a'; \theta_i^-) - Q(s,a; \theta_i) \right)^2 \right]$$



[Human-level control through deep reinforcement learning - V.Mnih et al. \(2014\)](#)



What about the instability problem?

Two solutions:

1. Experience replay

- store experience in a replay buffer
- randomly sample batches of data and train the network on them

2. Target-Value Network

- target-value network is a clone of our main neural network which we are training
- use it to calculate the target

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

- create a new clone once a while so that the target-value network is not too different from the trained network

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ε select a random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For

End For

Optimal Sepsis Treatment


Case study 1

Aniruddh Ragu et al. (2017)

Policy	Expected Return	Estimated Mortality
Physician	9.87	$13.9 \pm 0.5\%$
Normal Q-N	10.16	$12.8 \pm 0.5\%$
Autoencode Q-N	10.73	$11.2 \pm 0.4\%$

Optimal Sepsis Treatment

- training set: 15 500 survivors, 2 300 non-survivors
- physiological parameters including demographics, lab values and vital signs aggregated into windows of 4 hours (47 x 1 **state vector** for each patient at each timestep)
- 5 x 5 **action space** covering dosages of two different drugs
- continuous state space, discrete action space
- Deep Q-Network + Double Deep-Q + Dueling Q Network + Prioritized Experience Replay = Dueling Double-Deep Q Network (Dueling DDQN)
- evaluating models offline is difficult



[Ant Soccer \(DeepMind 2016\)](#)

Continuous action space

- vanilla Deep Q-Network cannot operate in continuous action space - discretize it?
- on the other hand, we can borrow some tricks introduced with DQN
- use a policy network - an Actor
- best performance when trained together with a Critic => **Actor-Critic**
- [Continuous Control with Deep Reinforcement Learning - T.P.Lillicrap et al. \(2015\)](#)

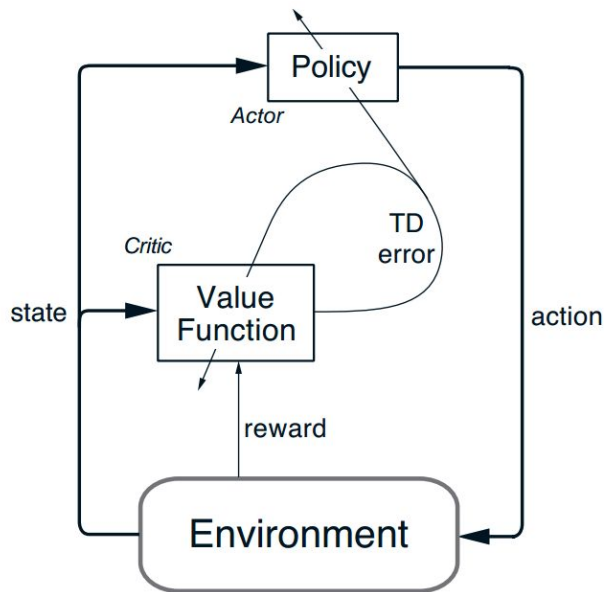


Figure 11.1: The actor-critic architecture.

Deep Deterministic Policy Gradients

Critic loss:

$$L(\theta^Q) = \mathbb{E}_{s_t \sim \rho^\beta, a_t \sim \beta, r_t \sim E} \left[(Q(s_t, a_t | \theta^Q) - y_t)^2 \right]$$

$$y_t = r(s_t, a_t) + \gamma Q(s_{t+1}, \mu(s_{t+1}) | \theta^Q).$$

Actor update:

$$\begin{aligned} \nabla_{\theta^\mu} J &\approx \mathbb{E}_{s_t \sim \rho^\beta} \left[\nabla_{\theta^\mu} Q(s, a | \theta^Q) \Big|_{s=s_t, a=\mu(s_t | \theta^\mu)} \right] \\ &= \mathbb{E}_{s_t \sim \rho^\beta} \left[\nabla_a Q(s, a | \theta^Q) \Big|_{s=s_t, a=\mu(s_t)} \nabla_{\theta^\mu} \mu(s | \theta^\mu) \Big|_{s=s_t} \right] \end{aligned}$$

- a gradient points in the direction of the steepest ascent of a function
- gradients of the Critic with respect to the action tell us how to improve policy network's parameters in order to maximize the Q-value

Algorithm 1 DDPG algorithm

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .

Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer R

for episode = 1, M **do**

 Initialize a random process \mathcal{N} for action exploration

 Receive initial observation state s_1

for $t = 1, T$ **do**

 Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise

 Execute action a_t and observe reward r_t and observe new state s_{t+1}

 Store transition (s_t, a_t, r_t, s_{t+1}) in R

 Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R

 Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$

 Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

 Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

 Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

end for

end for

You don't have to understand this algorithm completely. Just the basic idea behind actor-critic.

Cheetah

Low Dimensional Features

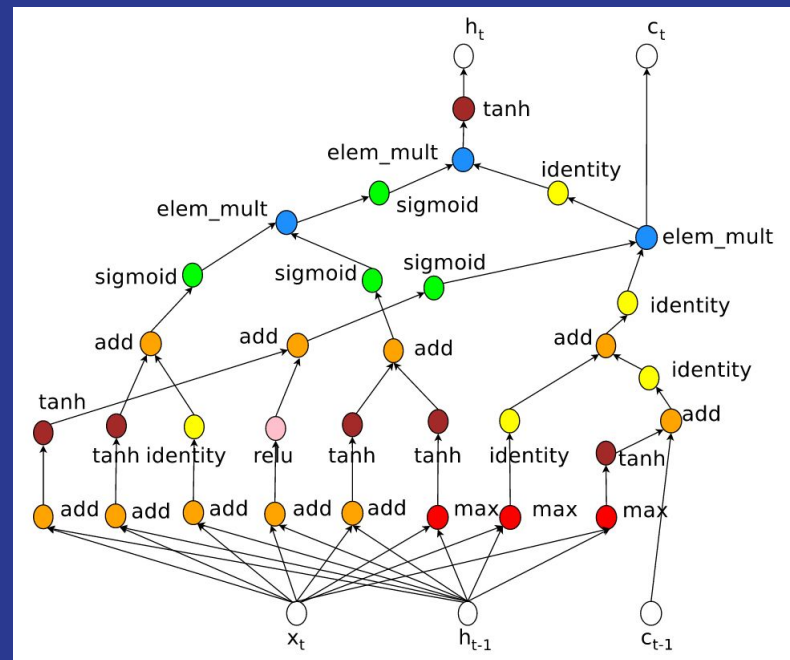


[DDPG demo](#)


Neural Architecture Search

Case study 2

Barret Zoph et al. (2017)




Convolutional Neural Networks




Source



Layer 1

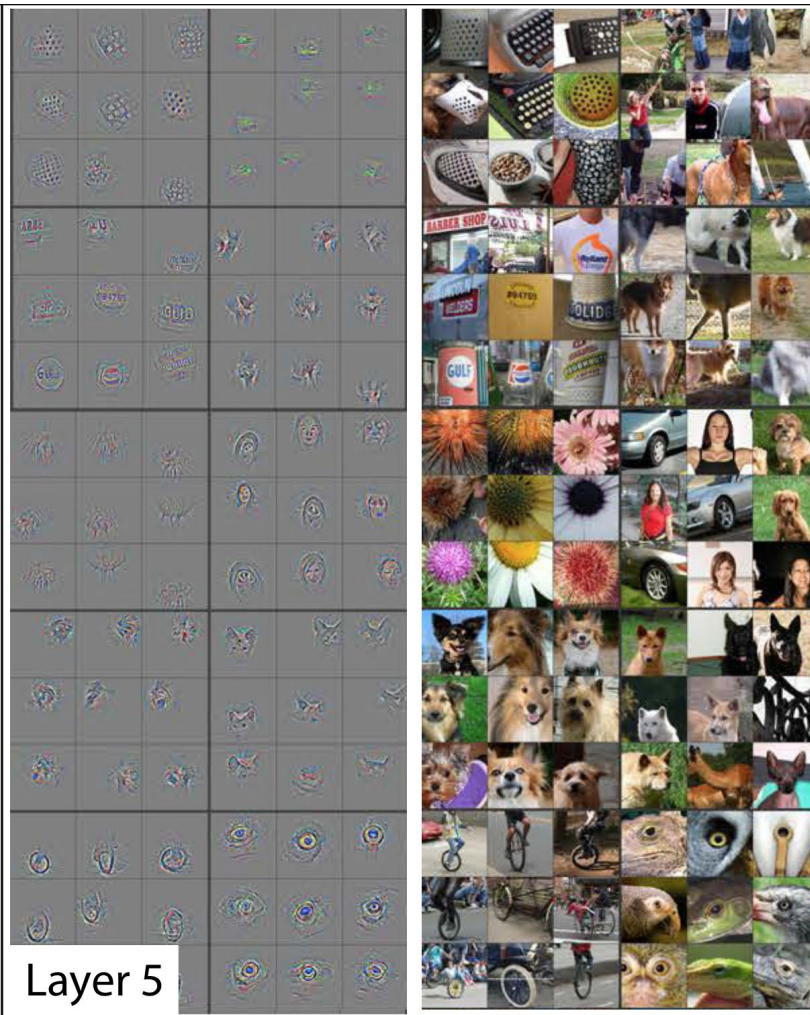
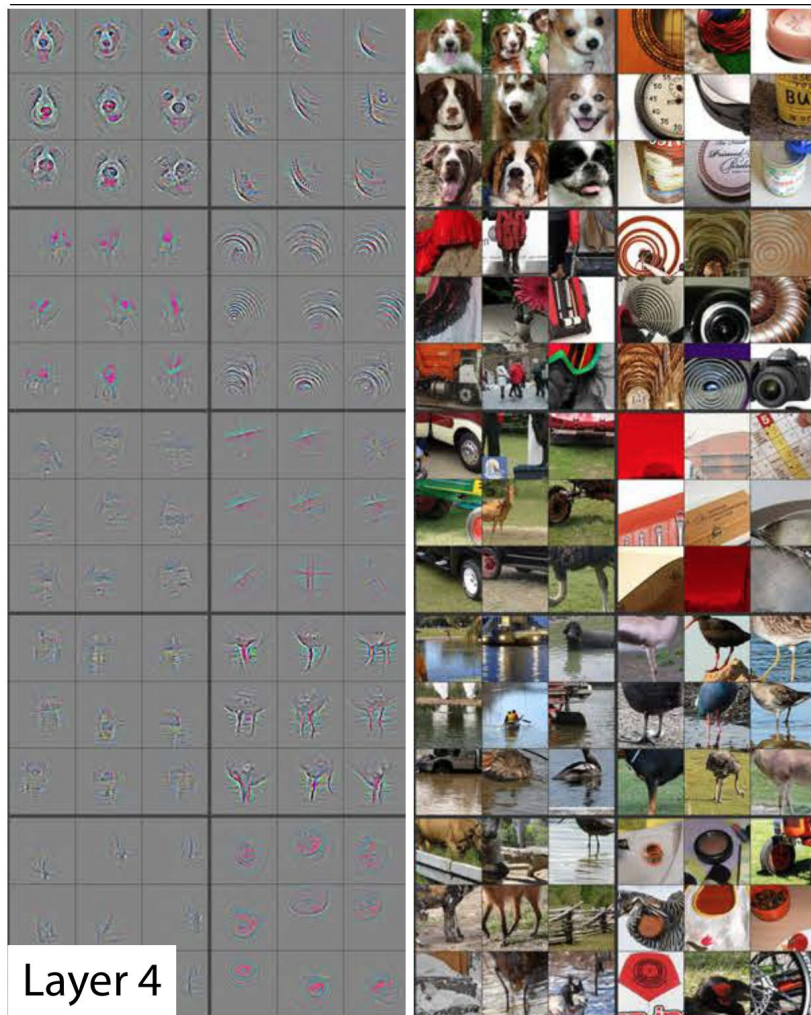


Layer 2



Layer 3





Neural Architecture Search

- train a RNN that generates hyperparameters for neural networks

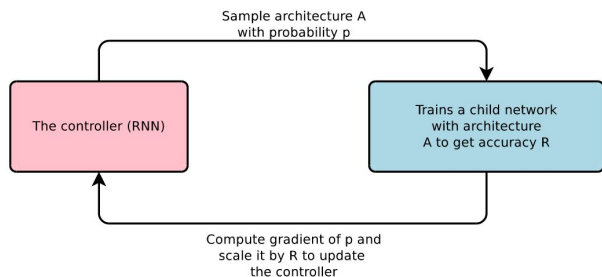


Figure 1: An overview of Neural Architecture Search.

- **reward**: the test accuracy of the generated network

- **update**: $\nabla_{\theta_c} J(\theta_c) = \sum_{t=1}^T E_{P(a_{1:T}; \theta_c)} [\nabla_{\theta_c} \log P(a_t | a_{(t-1):1}; \theta_c) R]$

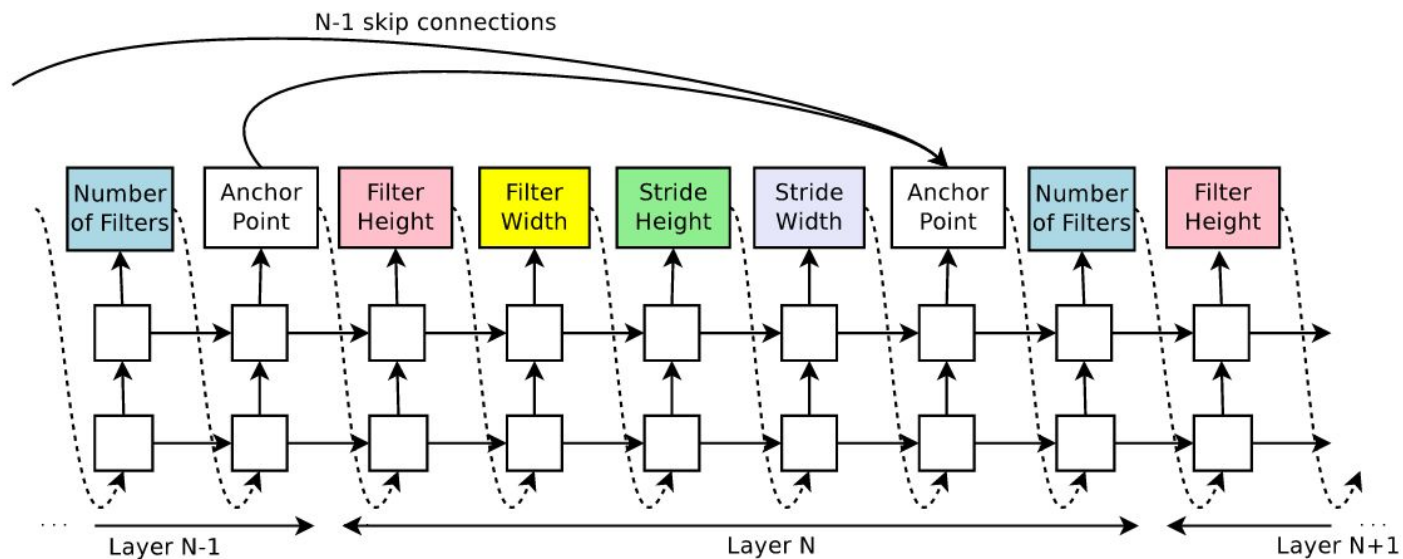


Figure 4: The controller uses anchor points, and set-selection attention to form skip connections.

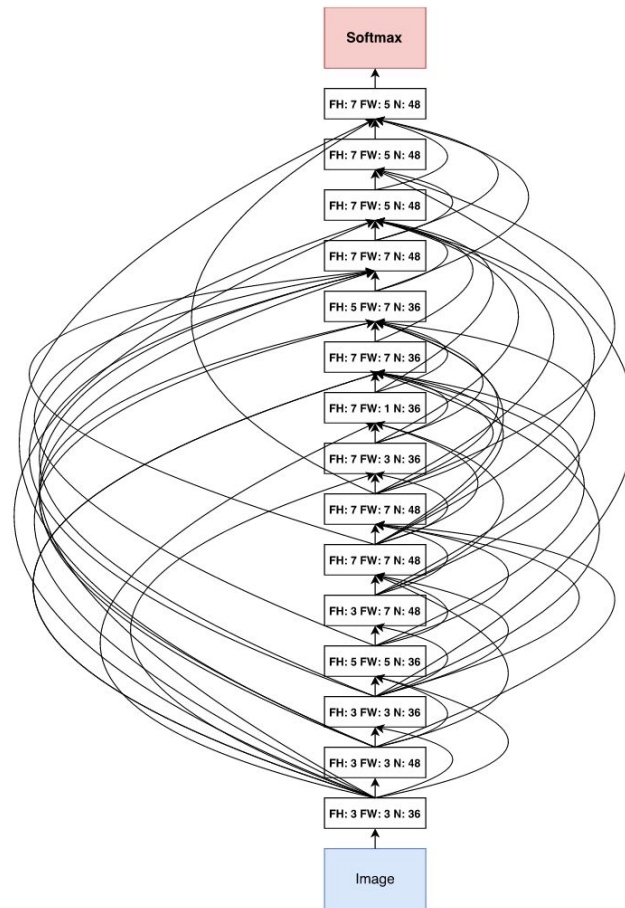


Figure 7: Convolutional architecture discovered by our method, when the search space does not have strides or pooling layers. FH is filter height, FW is filter width and N is number of filters. Note that the skip connections are not residual connections. If one layer has many input layers then all input layers are concatenated in the depth dimension.

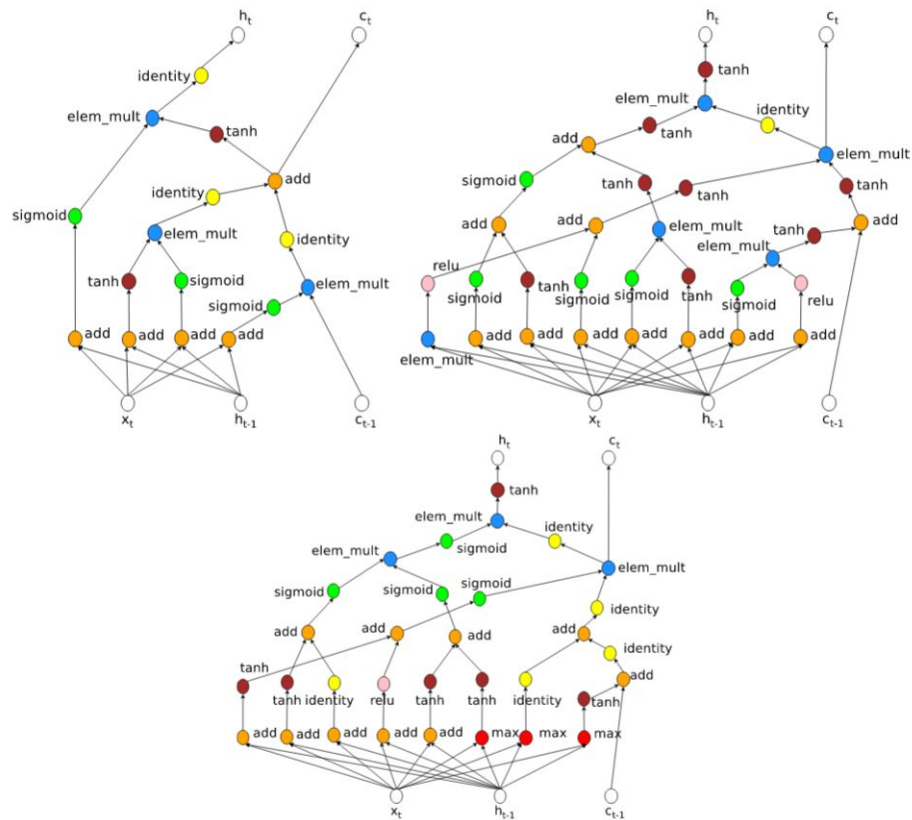


Figure 8: A comparison of the original LSTM cell vs. two good cells our model found. Top left: LSTM cell. Top right: Cell found by our model when the search space does not include \max and \sin . Bottom: Cell found by our model when the search space includes \max and \sin (the controller did not choose to use the \sin function).

Summary

- we've seen how to combine neural networks with reinforcement learning
- **Deep Q-Network** is one the first architectures that made this combination work
- for continuous states, we can use an Actor or a combination of an Actor and a Critic => Actor-Critic (for example the **DDPG** algorithm)
- there are many use cases of Deep Reinforcement Learning outside playing video games; on the other hand, video games are a great benchmark



The future of Deep Reinforcement Learning

- we've seen a lot of advancements in the Atari games environment but not many applications for real world problems; why?
- move towards general agents that are easy to deploy in many settings
- this could be accomplished by training agents in a modular fashion
- e.g. a module for vision, planning, motion and so on
- similar to how modern software is built

